

SpooofVane — User Guide

A task-based guide: for each thing you want to do, where to start, the steps, and what happens.

v0.5.0 · 29 May 2026 · Audience: SOC analysts, reviewers, brand admins, auditors, administrators

1. About this guide

This guide is organised by **task**, not by screen. Find what you want to accomplish in the contents below and jump to it. Each task tells you who can do it, what you need first, the exact steps (and which screen each step is on), the result, and what to do next. A short *Behind the scenes* note names the underlying capability if you want to go deeper.

Your role decides what you can do

Navigation and permitted actions depend on your role. From least to most privileged:

Viewer	Read dashboards and settings.
Auditor	+ Read the audit log and compliance evidence.
Reviewer	+ Approve or deny human-in-the-loop actions (must differ from the analyst).
Analyst	+ Triage alerts, investigate, request takedowns, onboard brands.
Admin	+ Manage users, agents, cost and tenants.
Owner	+ Manage tenants and billing.

The golden rule: every takedown and every sensitive action needs a second person's approval. You (the analyst) propose; a reviewer approves. This is by design and everything is recorded in a tamper-evident audit log.

2. Contents — what do you want to do?

GETTING STARTED

T1 Sign in and find your way around

SET UP PROTECTION

T2 Onboard a brand to monitor

T3 Protect an executive from deepfakes & impersonation

INVESTIGATE

T4 Triage the alert queue

T5 Review an alert's evidence and confirm the verdict

T6 Review a suspected deepfake

T7 See the campaign behind a single site

ACT

T8 Approve (or deny) a takedown — reviewer task

T9 Track takedown delivery

OVERSEE

T10 Search the audit trail

T11 Watch Bright Data spend against budget

ADMINISTER

T12 Manage agents and use the kill-switch

Getting started

T1 Sign in and find your way around

Purpose. Access SpoofVane and understand the layout before you begin work.

Who can do this: All roles

BEFORE YOU START

- A SpoofVane account from your administrator
- Your organisation's SSO, or an email + password
- An authenticator app if MFA is enabled

STEPS

1. **Open the console** Go to your SpoofVane URL. The **Login** screen (P01) appears.
2. **Choose a sign-in method** Click **Continue with SSO (OIDC)** or **SAML**, or enter email and password.
3. **Complete MFA if prompted** Enter the 6-digit code from your authenticator app.
4. **Orient yourself** You land on the **Dashboard** (P02). The left rail is your navigation; what you see there depends on your role.

✓ RESULT — WHAT HAPPENS

You are signed in on the Dashboard, showing today's KPI tiles, the threat-family breakdown, and the live processing pipeline.

→ WHAT TO DO NEXT

Most analysts go straight to **Triage queue** (T4). If you are setting up a new brand first, see **Onboard a brand** (T2).

BEHIND THE SCENES

Identity (G1) handles OIDC/SAML/MFA; RBAC (G2) decides which navigation items you can see.

Set up protection

T2 Onboard a brand to monitor

Purpose. Tell SpoofVane which brand to protect so discovery can begin finding impersonation.

Who can do this: Analyst, Admin, Owner

BEFORE YOU START

- You are signed in
- You know the brand's canonical login URL
- (Optional) a reference logo image

STEPS

1. **Open Brands** In the left rail, click **Brands** (P03). You see the list of brands already protected.
2. **Launch the wizard** Click + **New brand (wizard)**. The Brand Wizard opens.
3. **Enter the brand basics** Type the brand name and canonical login URL. The wizard *suggests* keywords, permutation TLDs and monitoring regions — accept or edit them.
4. **Set the escalation threshold** Choose the score threshold (default 0.65) above which a suspect is escalated for a verdict.
5. **Save** Click **Save**. You are taken to the **Brand detail** screen (P04) for the new brand.

✓ RESULT — WHAT HAPPENS

The brand is registered. Discovery sources immediately begin searching for impersonation; within minutes the Triage queue starts filling with suspects for this brand.

→ WHAT TO DO NEXT

Watch results arrive in the **Triage queue** (T4). To adjust assets or thresholds later, return to **Brand detail** (P04).

BEHIND THE SCENES

The Brand Wizard is module H3; saving the brand activates the Discovery group (A1-A10), which use Bright Data SERP, Residential Proxies, Datasets and more.

T3 Protect an executive from deepfakes & impersonation

Purpose. Monitor a named executive's likeness and voice across the web and social media.

Who can do this: Analyst (with a recorded lawful basis), Admin

BEFORE YOU START

- The brand is onboarded (T2)
- A recorded lawful basis for processing the executive's likeness/voice
- (Recommended) an enrolled reference voice sample

STEPS

1. **Open Exec protection** Click **Exec protection** (P09) in the left rail.
2. **Add the executive** Enter the executive's name and the social platforms to watch.
3. **Record the lawful basis** You will be required to enter a named lawful basis. **Until you do, monitoring will not start** — this is a deliberate privacy gate.
4. **Enrol a voiceprint (optional)** Upload a clean reference voice clip so cloned-voice media can be matched.
5. **Activate** Confirm. The exec surface monitor begins scanning.

✓ RESULT — WHAT HAPPENS

SpoofVane reports impersonating handles, likeness assets, and any cloned-voice or deepfake media targeting the executive, with a risk score.

→ WHAT TO DO NEXT

Suspect media appears in the **Deepfakes** queue (T6). Any takedown still flows through review (T8).

BEHIND THE SCENES

Exec attack surface is module H5; it is a 'sensitive agent' so module E8 (governance) enforces the lawful-basis gate (review §V8-5). Voiceprint matching is C9; deepfake scoring is C10.

Investigate

T4 Triage the alert queue

Purpose. Work through suspected impersonation sites and decide which need action.

Who can do this: Analyst

BEFORE YOU START

- At least one brand is onboarded and has produced suspects

STEPS

1. **Open Triage queue** Click **Triage queue** (P05). Each row is one suspect with a verdict pill, score, source and an SLA timer.
2. **Sort by urgency** Red **phish** pills and 'hot' SLA timers are most urgent. Start at the top.
3. **Open a suspect** Click a row to open the **Alert detail** screen (P06).

✓ RESULT — WHAT HAPPENS

You have the most urgent suspect open in Alert detail, ready to review the evidence.

→ WHAT TO DO NEXT

Review the evidence — see **Review an alert's evidence** (T5).

BEHIND THE SCENES

Suspects are scored by the Scoring group (C1–C8). The calibrated probability (C1) — not a raw sum — drives the verdict pill and ordering.

T5 Review an alert's evidence and confirm the verdict

Purpose. Decide whether a suspect is genuinely a phishing/impersonation site using the full evidence.

Who can do this: Analyst

BEFORE YOU START

- A suspect open in Alert detail (T4)

STEPS

1. **Read the Summary tab** The **Summary** tab shows the verdict, calibrated score, brand and when it was first seen.
2. **Inspect the visual & technical evidence** Open **Evidence** (screenshot/DOM), **Multi-region** (is it cloaking by country?), and **Kit** (which phishing kit matched, e.g. EvilProxy).
3. **Check the verdict trace** Open **Verdict trace** to see the model ensemble (Claude + GPT + Gemini). If the models disagreed, the confidence is discounted and the alert is flagged for human review.
4. **Review context tabs** **MITRE** shows mapped ATT&CK techniques; **Cluster** shows linked sites; **Deepfake** appears if media is involved.
5. **Decide** If it is malicious, click **Request takedown**. If not, mark it benign to dismiss.

✓ RESULT — WHAT HAPPENS

Your decision is recorded. A takedown request is queued for human approval; a dismissal closes the alert. Either way the action is written to the audit log.

→ WHAT TO DO NEXT

For a takedown, a reviewer approves it next — see **Approve a takedown** (T8). Track progress in **Takedowns** (T9).

BEHIND THE SCENES

Verdict ensemble is D1-D5; dissent handling is D5; MITRE enrichment is D7; clustering is C8. Your decision feeds the active-learning loop (D8/E6) that re-tunes calibration.

T6 Review a suspected deepfake

Purpose. Confirm or dismiss AI-generated video/voice impersonating your brand or executive.

Who can do this: Analyst

BEFORE YOU START

- A deepfake suspect exists (from exec protection T3 or discovery)

STEPS

1. **Open Deepfakes** Click **Deepfakes** (P08). Each item shows the multimodal score.
2. **Read the signal breakdown** Review face-embedding distance, lip-sync inconsistency, voiceprint match, and the C2PA provenance result. A valid provenance manifest *lowers* suspicion.
3. **Confirm the verdict** Accept the multimodal verdict (deepfake / suspicious / authentic) or override it.

✓ RESULT — WHAT HAPPENS

The deepfake is classified, with the per-signal reasoning retained for evidence.

→ WHAT TO DO NEXT

If action is needed, request a takedown; it flows through review (T8). A legal export carries a chain-of-custody pack.

BEHIND THE SCENES

Voiceprint is C9, deepfake ensemble is C10, multimodal fusion is D6, and the legal chain-of-custody pack is G10.

T7 See the campaign behind a single site

Purpose. Understand whether a suspect is part of a larger coordinated campaign.

Who can do this: Analyst

BEFORE YOU START

- At least two related suspects have been discovered

STEPS

1. **Open Clusters** Click **Clusters** (P07) to see the threat-network graph.
2. **Read the links** Nodes are connected when they share an ASN, registrar, phishing kit, favicon or TLS fingerprint.
3. **Jump to a node** Click any node to open that site's Alert detail.

✓ RESULT — WHAT HAPPENS

You can see the full campaign and act on it as a group rather than one site at a time.

→ WHAT TO DO NEXT

Take down the cluster's sites individually (each via T5 → T8).

BEHIND THE SCENES

Clustering is module C8 (union-find over shared signals); the graph is maintained incrementally by the cluster agent (E5).

Act

T8 Approve (or deny) a takedown — reviewer task

Purpose. Provide the mandatory human approval before SpoofVane submits a takedown.

Who can do this: Reviewer (must be a different person from the analyst who raised it)

BEFORE YOU START

- An analyst has requested a takedown (T5)
- You hold the Reviewer role

STEPS

1. **Open Review queue** Click **Review queue** (P12). Pending actions are listed.
2. **Examine the request** Review the suspect, its verdict and the drafted takedown packet.
3. **Approve or deny** Click **Approve** to authorise submission, or **Deny** with a reason.
4. **(Sensitive actions only)** If the action egresses synthetic credentials or a honeypot, a *second* authoriser is also required before it proceeds.

✓ RESULT — WHAT HAPPENS

On approval, the takedown is submitted to the correct registrar or hosting-abuse channel and a reference ID is returned. Every step is audited.

→ WHAT TO DO NEXT

Track delivery in **Takedowns** (T9).

BEHIND THE SCENES

HITL review is G4 (segregation of duties: reviewer ≠ analyst). The two-person rule for egress actions is enforced by governance (E8, review §V8-5).

T9 Track takedown delivery

Purpose. Monitor where each approved takedown is in the abuse-reporting pipeline.

Who can do this: Analyst, Reviewer

BEFORE YOU START

- A takedown has been approved (T8)

STEPS

1. **Open Takedowns** Click **Takedowns** (P10).
2. **Read the status** Each row shows the target URL, the channel (registrar or hosting provider), the status, and the reference ID.
3. **Override if needed** If a channel stalls, use the manual override to re-route or escalate.

✓ RESULT — WHAT HAPPENS

You have live visibility of every takedown and its reference for follow-up.

→ WHAT TO DO NEXT

Confirm resolution; the closed alert remains in **Audit** (T10) permanently.

BEHIND THE SCENES

Registrar adapters are F1-F3; hosting-abuse routing is F4; alerts also fan out to your SIEM/SOAR via F5 (Splunk/Sentinel/ServiceNow/PagerDuty/Cortex XSOAR).

Oversee

T10 Search the audit trail

Purpose. Find exactly who did what, when — for an investigation or an audit.

Who can do this: Auditor, Reviewer, Admin

BEFORE YOU START

- You hold a role with audit access

STEPS

1. **Open Audit** Click **Audit** (P11).
2. **Ask in plain language** Type a natural-language query such as '*analyst takedown actions for tenant acme*'. SpoofVane parses it into filters.
3. **Read the chained log** Each entry shows the actor, action, outcome and a hash linking it to the previous entry.

✓ RESULT — WHAT HAPPENS

You have a filtered, tamper-evident view of the relevant actions.

→ WHAT TO DO NEXT

For compliance reporting, continue to **Compliance** (P14) to export an auditor pack.

BEHIND THE SCENES

Natural-language search is H2; the log itself is the hash-chained audit (G3), which self-verifies and detects tampering.

T11 Watch Bright Data spend against budget

Purpose. Keep data-collection costs within the tenant's envelope.

Who can do this: Admin, Owner

BEFORE YOU START

- You hold the Admin or Owner role

STEPS

1. **Open Cost** Click **Cost** (P13).
2. **Read spend vs envelope** See current spend against the tier cap (free \$5 / pro \$100 / business \$1,000 / enterprise \$25,000) and the share handled by the cheap SLM path.
3. **Set or adjust alerts** Configure a budget alert threshold if needed.

✓ RESULT — WHAT HAPPENS

You can see and control data-collection cost before it exceeds the envelope.

→ WHAT TO DO NEXT

If spend is high, review which sources are most active in Brand detail (P04).

BEHIND THE SCENES

Cost tracking is G9; every Bright Data call records a cost row automatically. SLM-first routing (D4) keeps most low-severity verdicts on the cheap path.

Administer

T12 Manage agents and use the kill-switch

Purpose. Govern the autonomous agents and halt them instantly if anything looks wrong.

Who can do this: Admin

BEFORE YOU START

- You hold the Admin role

STEPS

1. **Open Admin · Agents** Click **Agents** (P15). You see each agent, its state and its blast-radius budget.
2. **Set governance budgets** Adjust the cost/risk envelope per agent if required.
3. **Use the kill-switch** To stop everything immediately, click **Global kill-switch**. To stop one tenant's agents only, use the per-tenant control.

✓ RESULT — WHAT HAPPENS

Agents run within their budgets; if you fire the kill-switch, every running agent halts and returns a halted state, recorded in the audit.

→ WHAT TO DO NEXT

Resume agents from the same screen when it is safe.

BEHIND THE SCENES

The kill-switch is E7 (global or per-tenant); governance is E8; every agent action is written to the hash-chained agent audit (E9).