

SpooofVane — User Guide

Every screen — shown with numbered steps, explained, and connected by the flows.

Product: SpooofVane · Coder: Claude (Opus 4.8) in Claude Code
Build: v0.5.0 – 29/05/2026 10:52 BST · all 21 console pages · Monzo demo seed

How this guide works

SpooofVane is a Security-Operations-Center (SOC) console for brand, executive, deepfake and fraud-surface defense. Each screen is shown with **numbered ● markers** placed on the elements you act on; the matching **numbered steps** tell you what to do. You also get **What it is · Key elements · Actions · Leads to →**. Part 2 is the end-to-end flows; Part 3 covers roles, page states and exports. Where a screen needs more than one page, the lower half is marked (*continued*).

Navigation map

GROUP	SCREEN	ROUTE
Operations	Dashboard	/
	Triage Queue	/triage
	Alert Detail	/alerts/:id
	Campaign Clusters	/clusters
	Deepfake Feed	/deepfakes
	Exec Protection	/exec-protection
	Brand Management	/brands
	Takedowns	/takedowns


	Cost Attribution	/cost
	Audit Log	/audit
	Review Queue	/review
Settings	API Keys / Webhooks / Integrations	/settings/*
Admin	Tenants	/admin/tenants
	Agent Control	/admin/agents
	Compliance	/admin/compliance
	Demo Health	/admin/demo-health
	Users (RBAC)	/admin/users
System	Not Found / Forbidden	404 · 403

Part 1 · Screen-by-screen

1 Sign in

/login

What it is. Access SpoofVane; the entry screen with SSO / SAML / local + MFA. Who: All roles



● The numbered markers on the screen match the numbered steps below.

1 Sign in (continued)

STEPS

- 1 Click SSO (OIDC) or SAML to use your identity provider.
- 2 Or type your email + password.
- 3 Click Sign in.
- 4 Enter the 6-digit MFA code if prompted — then you land on the Dashboard.

KEY ELEMENTS


- SSO (OIDC) and SAML buttons
- Email + password fallback
- TOTP MFA challenge when enabled

ACTIONS

Authenticate and reach the SOC console.

Leads to → Dashboard

What it is. Register and manage the brands SpoofVane protects; launch the onboarding wizard. **Who:** Analyst, Admin, Owner



● The numbered markers on the screen match the numbered steps below.

2 Brand Management *(continued)*

STEPS

- 1 Click + New brand (wizard) to start onboarding.
- 2 Enter the brand name and canonical login URL; accept the suggested keywords, TLDs and regions.
- 3 Pick the escalation score threshold (default 0.65), then save to open Brand detail.

KEY ELEMENTS

- Brand list with regions, canonical assets, detection thresholds
- Active discovery channels per brand
- Brand-onboarding wizard launcher

ACTIONS

Add a brand, launch the wizard, edit thresholds.

Leads to → Brand Detail · Dashboard (post-onboarding)

What it is. Monitor an executive's likeness and voice abuse across the open web (sensitive surface). Who: Analyst (lawful basis), Admin

The screenshot shows the SpoofVane v0.5 Enterprise interface. The sidebar on the left contains navigation links under 'MAIN' (Dashboard, Triage Queue, Campaign Clusters, Deepfake Feed, Exec Protection, Brand Management, Cost Attribution, Audit Log, Takedowns), 'SETTINGS' (API Keys, Webhooks, Integrations), and 'ADMIN' (Tenants, Agent Control, Compliance, System Health, Demo Health). The main content area is titled 'Executive Protection' and has a 'lawful-basis gated' status. It contains a form with five numbered steps: 1. Enter executive name and platforms, 2. Record a named lawful basis, 3. Optionally upload a reference voice clip, 4. Click Activate scanning, and 5. Review findings. The form is titled 'Executive Protection' and has a 'lawful-basis gated' status.

● The numbered markers on the screen match the numbered steps below.

3 Exec Protection *(continued)*

STEPS

- 1 Enter the executive's name and the platforms to watch.
- 2 Record a named lawful basis — required; monitoring will not start without it.
- 3 Optionally upload a reference voice clip for cloned-voice matching.
- 4 Click Activate scanning to begin.

KEY ELEMENTS


- Per-executive findings: likeness hits, voice exposure, PII leakage score
- Data-broker exposure + redaction task list
- Lawful-basis gate before any monitoring

ACTIONS

Review findings, trigger redaction, escalate (PII processing is gated).

Leads to → Alert Detail · Audit Log

What it is. The analyst's primary work surface — alerts awaiting a decision. Who: Analyst



- The numbered markers on the screen match the numbered steps below.

4 Triage Queue (continued)

STEPS

- 1 Scan the queue — each row is a suspect with its verdict pill, score, source and SLA timer.
- 2 Work top-down: red phish pills and short SLA timers are most urgent.
- 3 Click a row to open its Alert detail.

KEY ELEMENTS


- Virtualised table: severity, URL, family, score, SLA timer, status
- Filters, saved views, bulk-select
- CSV / NDJSON export

ACTIONS

Open, assign, change status, bulk-action, export.

Leads to → Alert Detail (click any row)

What it is. The full evidence dossier for one alert; where verdicts and takedowns are actioned. **Who:** Analyst



● The numbered markers on the screen match the numbered steps below.

5 Alert Detail *(continued)*

STEPS

- 1 Read the Summary tab — verdict, calibrated score, brand, first-seen.
- 2 Open Evidence + Multi-region (cloaking) tabs to see the captured page.
- 3 Open Verdict trace — the Claude+GPT+Gemini ensemble; dissent discounts confidence.
- 4 Check MITRE / Kit / Cluster tabs for context.
- 5 Click Request takedown (or mark benign to dismiss).

KEY ELEMENTS


- Header: family, threat flags, scoring breakdown, kit name/version
- Tabs: Summary, Evidence, Multi-region, Verdict, MITRE, Kit, Cluster, Deepfake, Takedown
- Every action writes a hash-chained audit entry

ACTIONS

Confirm / deny the verdict, draft & submit a takedown (HITL-gated), add notes.

Leads to → Campaign Clusters · Audit Log · Review Queue (on takedown)

What it is. Review queued deepfake media (image / audio / video) with model verdicts. Who: Analyst



- The numbered markers on the screen match the numbered steps below.

6 Deepfake Feed *(continued)*

STEPS

- 1 Open a queued item — it shows the multimodal score.
- 2 Read the per-signal breakdown; valid C2PA provenance lowers suspicion.
- 3 Confirm or override the multimodal verdict.

KEY ELEMENTS


- Face-embedding, lip-sync, voiceprint and C2PA-provenance scores
- Per-item verdict panel + chain-of-custody evidence pack

ACTIONS

Confirm / deny a verdict; open the chain-of-custody pack.

Leads to → Alert Detail (when tied to an alert)

What it is. See related threats as a network — kits / domains / hosting / actors in one campaign. Who: Analyst



● The numbered markers on the screen match the numbered steps below.

7 Campaign Clusters *(continued)*

STEPS

- 1 Open the threat-network graph.
- 2 Read the links — nodes connect by shared ASN, registrar, kit, favicon or TLS fingerprint.
- 3 Click any node to pivot into that site's Alert detail.

KEY ELEMENTS


- Interactive graph (shared ASN, kit family, registrar, favicon, TLS)
- Cluster-member side panel with attacker profile

ACTIONS

Select a node; pivot to a member's Alert Detail.

Leads to → Alert Detail

What it is. The human-in-the-loop approval surface; reviewer ≠ analyst is enforced.
Who: Reviewer (≠ analyst)



● The numbered markers on the screen match the numbered steps below.

8

Review Queue *(continued)*

STEPS

- 1 Open the queue — pending actions are listed.
- 2 Examine the suspect, verdict and drafted packet.
- 3 Click Approve to authorise (or Deny with a reason).
- 4 For egress actions (canary creds / honeypot), add a second authoriser.

KEY ELEMENTS

- Pending takedowns/verdicts with the drafted packet + verdict trace
- Second-authoriser step for egress actions

ACTIONS

Approve / deny; provide a second authorisation where required.

Leads to → Takedowns (on approval) · Audit Log

What it is. Multi-channel takedown status across registrars and hosting-abuse channels. **Who:** Analyst, Reviewer

SpoofVane
v0.5 Enterprise

Search alerts, URLs, domains...

Monzo Bank ▾

Takedown status

1 URL	CHANNEL	STATUS	REF
acme-secure.top	Cloudflare	2 submitted	CF-6d8f
my-acme.cc	3 Hetzner abuse	HITL pending	-
acme-verify.xyz	Namecheap	acknowledged	NC-44a1

● The numbered markers on the screen match the numbered steps below.

9 Takedowns (continued)

STEPS

- 1 Open Takedowns — each row shows URL, channel, status and reference ID.
- 2 Read the status — registrar or hosting-abuse channel and its state.
- 3 Re-route or escalate a stalled takedown manually.

KEY ELEMENTS


- Per-takedown channel, status, reference ID
- Manual override on stalled channels

ACTIONS

Track delivery; override a stalled channel.

Leads to → Audit Log

What it is. The tamper-evident record of every consequential action. Who: Auditor, Reviewer, Admin



● The numbered markers on the screen match the numbered steps below.

10 Audit Log *(continued)*

STEPS

- 1 Open the audit log — the hash-chained trail.
- 2 Type a natural-language query, e.g. 'analyst takedown actions for tenant acme'.
- 3 Read results — actor, action, outcome and the hash linking each entry.

KEY ELEMENTS


- Hash-chained trail (Blake2b + Ed25519), 7-year retention
- Natural-language search + chain verification

ACTIONS

Search, verify the chain, export the audit pack.

Leads to → Alert Detail (each entry links to its subject)

What it is. Per-tenant spend on Bright Data products and LLM / SLM calls. Who: Admin, Owner



● The numbered markers on the screen match the numbered steps below.

11 Cost Attribution *(continued)*

STEPS

- 1 Open Cost — see spend against the tier envelope.
- 2 Read the bar: current spend, the cap, and the SLM cheap-path share.
- 3 Set a budget alert threshold if needed.

KEY ELEMENTS


- Spend by product / action / tenant, budget bars
- Bright Data rows tagged source=brd_*; SLM-first routing share

ACTIONS

Set budgets; export the cost-proof report.

Leads to → Compliance · Audit Log

What it is. The control room for the autonomous agents — budgets, governance, kill-switch. **Who:** Admin



- The numbered markers on the screen match the numbered steps below.

12 Agent Control *(continued)*

STEPS

- 1 Review the agent table — each agent, its state and blast-radius budget.
- 2 Adjust the cost / risk envelope per agent.
- 3 Use Global kill-switch to halt everything (or scope it per-tenant).

KEY ELEMENTS

- Agent registry with status + governance budgets
- Global Kill-Switch (≤ 5 s)
- Governance banner (NIST AI RMF + EU AI Act + UK ECCTA)

ACTIONS

Start / pause agents, set envelopes, halt all immediately.

Leads to → Audit Log (every action recorded)

Part 2 • Screen flows

A — Daily triage (the core loop)

Dashboard → Triage Queue → Alert Detail → Evidence/Verdict tabs → Request takedown → Review Queue (HITL) → Takedowns → Audit Log → back to the queue.

B — Brand onboarding

Brand Management → wizard → brand + regions + thresholds → discovery (A1-A10) provisions → Dashboard populates → Brand Detail.

C — Deepfake review

Dashboard Deepfakes KPI → Deepfake Feed → open a clip → confirm/deny → chain-of-custody pack → Alert Detail.

D — Executive protection

Exec Protection → pick executive → record lawful basis → review findings → redaction → escalate → Alert Detail / Audit.

E — Campaign investigation

Alert Detail → Cluster tab → Campaign Clusters graph → pivot across kit/hosting/actor → related Alert Details.

F — Agent governance & kill-switch

Agent Control → inspect status + budgets → (incident) Global Kill-Switch → all agents halt ≤ 5 s → Audit Log entry.

G — Cost & compliance

Cost Attribution → spend + budgets → Compliance → framework status → export signed evidence pack.

H — Audit investigation

Audit Log → NL search → open a result → Alert Detail → verify hash-chain integrity.

Part 3 · Roles, states & exports

Roles & visibility

Viewer	Read dashboards and settings.
Auditor	+ Read audit & compliance evidence.
Reviewer	+ Approve / deny HITL actions; reviewer ≠ analyst enforced.
Analyst	+ Triage, investigate, request takedowns, onboard brands.
Admin	+ Manage users, agents, cost, tenants, integrations.
Owner	+ Manage tenants and billing.

Forbidden actions are hidden; a forbidden deep link shows the **Forbidden (403)** screen.

Page states & exports. Every data screen renders one of four states — **Populated**, **Empty**, **Loading** (skeletons), **Error** (with retry). Tables export CSV / NDJSON; alerts export a signed PDF; Compliance exports a signed evidence pack; deepfake evidence exports a chain-of-custody bundle.