



SpoofVane — User Guide

Every screen — shown with numbered steps, explained, and connected by the flows.

Product: SpoofVane · Coder: Claude (Opus 4.8) in Claude Code
Build: v0.5.0 — 29/05/2026 11:20 BST · all 21 console pages · Monzo demo seed

How this guide works

SpoofVane is a Security-Operations-Center (SOC) console for brand, executive, deepfake and fraud-surface defense. For each screen you get a **capture with numbered ● markers** beside the matching **numbered steps**, plus **What it is · Key elements · Actions · Leads to →**. Part 2 is the end-to-end flows; Part 3 covers roles, page states and exports. Each screen is kept together on one page; if a screen ever runs onto the next page, only its header is repeated, marked (*continued*).

Navigation map

GROUP	SCREEN	ROUTE
Operations	Dashboard	/
	Triage Queue	/triage
	Alert Detail	/alerts/:id
	Campaign Clusters	/clusters
	Deepfake Feed	/deepfakes
	Exec Protection	/exec-protection
	Brand Management	/brands
	Takedowns	/takedowns
	Cost Attribution	/cost
	Audit Log	/audit
	Review Queue	/review
Settings	API Keys / Webhooks / Integrations	/settings/*
Admin	Tenants	/admin/tenants

	Agent Control	/admin/agents
	Compliance	/admin/compliance
	Demo Health	/admin/demo-health
	Users (RBAC)	/admin/users
System	Not Found / Forbidden	404 · 403

Part 1 · Screen-by-screen

1

Dashboard

/

What it is. The SOC at-a-glance landing screen after sign-in. Who: All roles



● Numbered markers match the numbered steps.

KEY ELEMENTS

- KPI tiles: Total Alerts, Open Alerts, Deepfakes (24h), Active Clusters, Exec Protection, Takedown Rate
- Alert-trend chart, Top Targeted Brands, Live Pipeline strip, Attack-Family distribution
- Capability Status (agents) and Compliance Health tiles

STEPS


- 1 Read the KPI tiles for today's posture.
- 2 Scan the trend chart and top targeted brands.
- 3 Click any KPI tile or chart segment to open the filtered view behind it.

ACTIONS

Click a KPI tile or chart segment to jump into the filtered view.

Leads to → Triage Queue · Deepfake Feed · Campaign Clusters · Exec Protection · Cost Attribution

What it is. The analyst's primary work surface — alerts awaiting a decision. **Who:** Analyst



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Virtualised alert table: severity, brand, suspect URL, family, score, SLA timer, status
- Filters, saved views, bulk-select
- CSV / NDJSON export

STEPS


- 1 Scan the queue — each row is a suspect with verdict pill, score, source and SLA timer.
- 2 Work top-down: red phish pills and short SLA timers are most urgent.
- 3 Click a row to open its Alert detail.

ACTIONS

Open an alert, assign, change status, bulk-action, export.

Leads to → Alert Detail (click any row)

What it is. The full evidence dossier for one alert; where verdicts and takedowns are actioned. Who: Analyst



● Numbered markers match the numbered steps.

STEPS

- 1 Read the Summary tab — verdict, calibrated score, brand, first-seen.
- 2 Open Evidence + Multi-region tabs to see the captured page and check cloaking.
- 3 Open Verdict trace — the Claude+GPT+Gemini ensemble; dissent discounts confidence.
- 4 Check MITRE / Kit / Cluster tabs for context.
- 5 Click Request takedown, or mark benign to dismiss.

KEY ELEMENTS

- Header: family, assignee, threat flags, scoring breakdown, kit name/version/confidence
- Tabs: Summary, Evidence, Verdict, Multi-region (cloaking), MITRE, Kit, Cluster, Deepfake, Takedown
- Every consequential action writes a hash-chained audit entry

ACTIONS

Confirm / deny the verdict, draft & submit a takedown (HITL-gated), add notes.

Leads to → Campaign Clusters (Cluster tab) · Audit Log (Ledger tab) · Review Queue (on takedown)

What it is. See related threats as a network — kits / domains / hosting / actors in one campaign. Who: Analyst



● Numbered markers match the numbered steps.

STEPS

- 1 Open the threat-network graph.
- 2 Read the links — nodes connect by shared ASN, registrar, kit, favicon or TLS.
- 3 Click any node to pivot into that site's Alert detail.

KEY ELEMENTS


- Interactive threat-network graph (shared hosting ASN, kit family, registrar, favicon, TLS fingerprint)
- Cluster-member side panel with attacker profile and recommended actions
- STIX export of the whole cluster

ACTIONS

Select a node to inspect; pivot from a member into its Alert Detail.

Leads to → Alert Detail

What it is. Review queued deepfake media (image / audio / video) with model verdicts. **Who:** Analyst



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Queue of clips/images with face-embedding, lip-sync, voiceprint and C2PA-provenance scores
- Per-item verdict panel and evidence pack
- Chain-of-custody export for legal discovery

STEPS


- 1 Open a queued item — it shows the multimodal score.
- 2 Read the per-signal breakdown; valid C2PA provenance lowers suspicion.
- 3 Confirm or override the multimodal verdict; open the chain-of-custody pack if needed.

ACTIONS

Confirm / deny a deepfake verdict; open the chain-of-custody pack.

Leads to → Alert Detail (when tied to an alert)

What it is. Monitor executive likeness and voice abuse across the open web (sensitive surface). Who: Analyst (lawful basis), Admin



● Numbered markers match the numbered steps.

STEPS

- 1 Enter the executive's name and the platforms to watch.
- 2 Record a named lawful basis — required; monitoring will not start without it.
- 3 Optionally upload a reference voice clip for cloned-voice matching.
- 4 Click Activate scanning to begin.

KEY ELEMENTS


- Per-executive surface findings (likeness hits, voice exposure from OSINT crawlers)
- PII leakage score, data-broker exposure, redaction task list
- Lawful-basis gate before any monitoring starts

ACTIONS

Review a finding, trigger a redaction recommendation, escalate (PII processing is gated).

Leads to → [Alert Detail](#) · [Audit Log](#)

What it is. Manage the brands SpoofVane protects; launch the onboarding wizard. **Who:** Analyst, Admin, Owner



● Numbered markers match the numbered steps.

STEPS

- 1 Click + New brand (wizard) to start onboarding.
- 2 Enter the brand name and canonical login URL; accept the suggested keywords, TLDs and regions.
- 3 Pick the escalation score threshold (default 0.65), then save to open Brand detail.

KEY ELEMENTS


- Brand list with protected regions, canonical assets and detection thresholds
- Active discovery channels per brand
- Brand-onboarding wizard launcher

ACTIONS

Add a brand, launch the wizard, edit thresholds.

Leads to → Brand Detail · Dashboard (post-onboarding)

What it is. Per-tenant spend on Bright Data products and LLM / SLM calls. **Who:** Admin, Owner



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Spend breakdown by product / action / tenant, budget bars, alerting thresholds
- Bright Data rows tagged source=brd_*; SLM-first routing share
- Cost-proof export (rules in docs/COST_ENVELOPE.md)

STEPS


- 1 Open Cost — see spend against the tier envelope.
- 2 Read the bar: current spend, the cap, and the SLM cheap-path share.
- 3 Set a budget alert threshold if needed.

ACTIONS

Set budgets; export the cost-proof report.

Leads to → Compliance · Audit Log

What it is. The tamper-evident record of every consequential action. **Who:** Auditor, Reviewer, Admin



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Searchable hash-chained trail (Blake2b + Ed25519 signed), 7-year retention
- Natural-language search
- Chain verification + audit-pack export

STEPS


- 1 Open the audit log — the hash-chained trail.
- 2 Type a natural-language query, e.g. 'analyst takedown actions for tenant acme'.
- 3 Read results — actor, action, outcome and the hash linking each entry.

ACTIONS

Search, verify the chain, export the audit pack.

Leads to → Alert Detail (each entry links back to its subject)

What it is. The control room for the autonomous agents (J1-J10). Who: Admin



● Numbered markers match the numbered steps.

STEPS

- 1 Review the agent table — each agent, its state and blast-radius budget.
- 2 Adjust the cost / risk envelope per agent.
- 3 Use Global kill-switch to halt everything (or scope it per-tenant).

KEY ELEMENTS


- Agent registry with status and governance budgets
- Global Kill-Switch (propagates in ≤ 5 s)
- Governance framework banner (NIST AI RMF + EU AI Act + UK ECCTA)

ACTIONS

Start / pause an agent, set budget / risk envelopes, stop every agent immediately.

Leads to → Audit Log (every action recorded)

What it is. Compliance evidence at a glance. Who: Auditor, Admin



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Framework tiles: SOC 2, ISO 27001, GDPR, NIST AI RMF, EU AI Act, ECCTA
- Per-framework control mappings + evidence-pack export
- Board declaration pack for presentation

STEPS


- 1 Read the framework tiles and average score.
- 2 Drill into a framework to see its control mappings and evidence.
- 3 Export a signed compliance / board pack.

ACTIONS

Export a signed compliance pack;
drill into a control's evidence.

Leads to → The compliance/ document pack

What it is. Multi-tenant administration. Who: Owner, Admin



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Tenant list with tier, region, status, isolation settings
- Per-tenant monthly cost and alert volume
- Data-residency region per tenant

STEPS


- 1 Review the tenant list — tier, alerts, cost, status.
- 2 Set the data-residency region for a tenant.
- 3 Create or suspend a tenant as needed.

ACTIONS

Create / suspend a tenant, set data-residency region.

Leads to → Users (RBAC)

What it is. Internal readiness check — confirms the seed populates every page and acceptance criteria pass. Who: Admin



● Numbered markers match the numbered steps.

KEY ELEMENTS

- Green-tick matrix of pages × seed completeness
- 41-point acceptance-criteria coverage
- Per-threshold required vs actual counts

STEPS

- 1 Read the ALL PASS banner and pass-rate.
- 2 Review the seed counts — required vs actual per threshold.
- 3 Re-run the seed check; jump to any under-seeded page.


ACTIONS

Re-run the seed check; jump to any under-seeded page.

Leads to → Any under-seeded page

What it is. Multi-channel takedown status across registrars and hosting-abuse channels.

Who: Analyst, Reviewer



● Numbered markers match the numbered steps.

STEPS

- 1 Open Takedowns — each row shows URL, channel, status and reference ID.
- 2 Read the status — registrar or hosting-abuse channel and its state.
- 3 Re-route or escalate a stalled takedown manually.

KEY ELEMENTS


- Per-takedown channel (Cloudflare / GoDaddy / Namecheap / hosting abuse), status, reference ID
- Manual override on stalled channels
- Audit trail for every submission

ACTIONS

Track delivery; re-route or escalate a stalled takedown.

Leads to → Audit Log

What it is. The human-in-the-loop approval surface; reviewer $\neq$ analyst is enforced. Who: Reviewer ($\neq$ analyst)



● Numbered markers match the numbered steps.

STEPS

- 1 Open the queue — pending actions are listed.
- 2 Examine the suspect, verdict and drafted packet.
- 3 Click Approve to authorise (or Deny with a reason).
- 4 For egress actions, add a second authoriser.

KEY ELEMENTS


- Pending takedowns/verdicts with the drafted packet and full verdict trace
- Second-authoriser step for egress actions (canary creds / honeypot)
- Approve / deny with reason, all audited

ACTIONS

Approve or deny; provide a second authorisation where required.

Leads to → Takedowns (on approval) · Audit Log

What it is. The entry screen — SSO / SAML / local credentials with MFA. **Who:** All roles



● Numbered markers match the numbered steps.

STEPS

- 1 Click SSO (OIDC) or SAML to use your identity provider.
- 2 Or type your email + password.
- 3 Click Sign in.
- 4 Enter the 6-digit MFA code if prompted — then you land on the Dashboard.

KEY ELEMENTS

- SSO (OIDC) and SAML buttons
- Email + password fallback
- TOTP MFA challenge when enabled

ACTIONS

Authenticate and reach the SOC console.

Leads to → Dashboard

Part 2 • Screen flows

A — Daily triage (the core loop)

Dashboard → Open Alerts KPI → Triage Queue (filter, pick an alert) → Alert Detail → Evidence tab (screenshots, DOM, multi-region cloaking) → Verdict tab (confirm/deny) → Takedown tab (draft) → Review Queue (HITL approve) → Takedowns (track) → Audit Log (hash-chained) → back to the queue.

B — Brand onboarding

Brand Management → Add brand / launch wizard → enter brand + regions + thresholds → wizard provisions discovery (A1-A10) → Dashboard begins to populate alerts → Brand Detail.

C — Deepfake review

Dashboard Deepfakes KPI → Deepfake Feed → open a clip (face / voiceprint / lip-sync / C2PA scores) → confirm/deny → chain-of-custody pack → Alert Detail (if tied to an alert).

D — Executive protection

Exec Protection → pick an executive → review likeness / voice findings → record lawful basis → action a redaction recommendation → escalate → Alert Detail / Audit Log.

E — Campaign investigation

Triage / Dashboard → Alert Detail → Cluster tab → Campaign Clusters graph → pivot across kit / hosting / actor nodes → open related Alert Details.

F — Agent governance & kill-switch

Admin › Agent Control → inspect agent status + budgets → (incident) press the Global Kill-Switch → all agents halt ≤ 5 s → Audit Log confirms with a hash-chained entry.

G — Cost & compliance review

Cost Attribution → review per-tenant Bright Data + LLM spend, set budgets → Admin › Compliance → check framework status → export signed evidence pack.

H — Audit investigation

Audit Log → NL search (“reviewer overrides last week”) → open a result → Alert Detail (the subject) → verify hash-chain integrity.

Part 3 · Roles, states & exports

Roles & visibility

Viewer	Read dashboards and settings.
Auditor	+ Read the audit log and compliance evidence.
Reviewer	+ Approve / deny HITL actions; reviewer ≠ analyst is enforced.
Analyst	+ Triage, investigate, request takedowns, onboard brands.
Admin	+ Manage users, agents, cost, tenants, integrations.
Owner	+ Manage tenants and billing.

Forbidden actions are hidden; a forbidden deep link shows the **Forbidden (403)** screen.

Page states & exports. Every data screen renders one of four states — **Populated**, **Empty**, **Loading** (skeletons), **Error** (with retry). Tables export CSV / NDJSON; alerts export a signed PDF; Compliance exports a signed evidence pack; deepfake evidence exports a chain-of-custody bundle.

SpoofVane User Guide · Claude (Opus 4.8) in Claude Code · 29/05/2026 11:20 BST · screens P01–P21
per console/src/lib/pages.ts · capabilities A–H per docs/TRACEABILITY_MASTER.md. Captures are
schematic renders in the product's real design tokens.