

SpooofVane — User Guide

Every screen — shown, explained, and connected by the end-to-end flows.

Product: SpooofVane · Coder: Claude (Opus 4.8) in Claude Code
Build: v0.5.0 – 29/05/2026 10:32 BST · all 21 console pages · Monzo demo seed

How this guide works

SpooofVane is a Security-Operations-Center (SOC) console for brand, executive, deepfake and fraud-surface defense. Every authenticated screen renders inside a fixed left sidebar (grouped into **Operations**, **Settings** and **Admin**) and a scrollable content pane, in a dark-first SOC theme. For each screen this guide gives you **What it is** · **Key elements** · **Actions** · **Leads to** →, beside a capture of the screen. Part 2 is the screen-by-screen tour; Part 3 is the end-to-end flows; Part 4 covers roles, page states and exports.


Navigation map

GROUP	SCREEN	ROUTE
Operations	Dashboard	/
	Triage Queue	/triage
	Alert Detail (drill-down)	/alerts/:id
	Campaign Clusters	/clusters
	Deepfake Feed	/deepfakes
	Exec Protection	/exec-protection
	Brand Management	/brands
	Brand Detail	/brands/:id

	Takedowns	/takedowns
	Cost Attribution	/cost
	Audit Log	/audit
	Review Queue	/review
Settings	API Keys	/settings/api-keys
	Webhooks	/settings/webhooks
	Integrations	/settings/integrations
	Settings	/settings
Admin	Tenants	/admin/tenants
	Agent Control	/admin/agents
	Compliance	/admin/compliance
	Demo Health	/admin/demo-health
	Users (RBAC)	/admin/users
System	Not Found / Forbidden	404 · 403

Part 1 · Screen-by-screen

What it is. The SOC at-a-glance landing screen after sign-in.



KEY ELEMENTS

- KPI tiles: Total Alerts, Open Alerts, Deepfakes (24h), Active Clusters, Exec Protection, Takedown Rate
- Alert-trend chart, Top Targeted Brands, Live Pipeline strip, Attack-Family distribution
- Capability Status (agents) and Compliance Health tiles

ACTIONS

Click a KPI tile or chart segment to jump into the filtered view behind it.

Leads to → Triage Queue · Deepfake Feed · Campaign Clusters · Exec Protection · Cost Attribution

What it is. The analyst's primary work surface — alerts awaiting a decision.

SpoofVane
v0.5 Enterprise

Search alerts, URLs, domains...

3 Monzo Bank ▼

Triage Queue
83 alerts · None selected

Search URL or family... All severities ▼ All statuses ▼

SEVERITY STATUS URL FAMILY REGION DF CLUSTER SCORE

critical	new	https://monzo-secure-login.co.uk/auth/0	banking	DE	92	c1-002	
high	in progress	https://monzobank-verify.com/auth/1	M365	US		c1-003	
medium	true positive	https://login-monzo.net/auth/2	crypto	SG		c1-004	
low	false positive	https://monzo-account.com/auth/3	payment	AU	84	c1-005	
critical	takedown	https://secure-monzo.online/auth/4	support	GB		c1-006	

SpoofVane v0.5 · 21 routes · 10 agents

KEY ELEMENTS


- Virtualised alert table: severity, brand, suspect URL, family, score, SLA timer, status
- Filters, saved views, bulk-select
- CSV / NDJSON export

ACTIONS

Open an alert, assign it, change status, bulk-action, export the queue.

Leads to → Alert Detail (click any row)

What it is. The full evidence dossier for one alert; where verdicts and takedowns are actioned.



KEY ELEMENTS


- Header: family, assignee, threat flags, scoring breakdown, kit name/version/confidence
- 8 evidence tabs: Overview, Evidence, Verdict, Multi-region (cloaking), MITRE, Ledger, Cluster, Takedown
- Every consequential action writes a hash-chained audit entry

ACTIONS

Confirm / deny the verdict, draft & submit a takedown (behind a human-in-the-loop gate), add notes.

Leads to → Campaign Clusters (Cluster tab) · Audit Log (Ledger tab) · Review Queue (on takedown)

What it is. See related threats as a network — kits / domains / hosting / actors in one campaign.



KEY ELEMENTS


- Interactive threat-network graph (shared hosting ASN, kit family, registrar, favicon, TLS fingerprint)
- Cluster-member side panel with attacker profile and recommended actions

ACTIONS

Select a node to inspect; pivot from a member into its Alert Detail.

Leads to → Alert Detail

What it is. Review queued deepfake media (image / audio / video) with model verdicts.



KEY ELEMENTS


- Queue of clips/images with face-embedding, lip-sync, voiceprint and C2PA-provenance scores
- Per-item verdict panel and evidence pack

ACTIONS

Confirm / deny a deepfake verdict; open the chain-of-custody pack.

Leads to → Alert Detail (when tied to an alert)

What it is. Monitor executive likeness and voice abuse across the open web.



KEY ELEMENTS


- Per-executive surface findings (likeness hits, voice exposure from OSINT crawlers)
- PII leakage score, data-broker exposure, redaction task list

ACTIONS

Review a finding, trigger a redaction recommendation, escalate (PII processing is lawful-basis gated).

Leads to → [Alert Detail](#) · [Audit Log](#)

What it is. Manage the brands SpoofVane protects.



KEY ELEMENTS


- Brand list with protected regions, canonical assets and detection thresholds
- Active discovery channels per brand
- Brand-onboarding wizard launcher

ACTIONS

Add a brand, launch the wizard, edit thresholds.

Leads to → Brand Detail · Dashboard (post-onboarding)

What it is. Per-tenant spend on Bright Data products and LLM / SLM calls.



KEY ELEMENTS


- Spend breakdown by product / action / tenant, budget bars, alerting thresholds
- Bright Data rows tagged source=brd_*; SLM-first routing share

ACTIONS

Set budgets; export the cost-proof report (rules in docs/COST_ENVELOPE.md).

Leads to → Compliance · Audit Log

What it is. The tamper-evident record of every consequential action.



KEY ELEMENTS


- Searchable hash-chained trail (Blake2b + Ed25519 signed), 7-year retention
- Natural-language search
- Chain verification + audit-pack export

ACTIONS

Search, verify the chain, export the audit pack.

Leads to → Alert Detail (each entry links back to its subject)

What it is. The control room for the autonomous agents (J1-J10).



KEY ELEMENTS


- Agent registry with status and governance budgets
- Global Kill-Switch (propagates in ≤ 5 s)
- Governance framework banner (NIST AI RMF + EU AI Act + UK ECCTA)

ACTIONS

Start / pause an agent, set budget / risk envelopes, stop every agent immediately.

Leads to → Audit Log (every action recorded)

What it is. Compliance evidence at a glance.



KEY ELEMENTS


- Framework tiles: SOC 2, ISO 27001, GDPR, NIST AI RMF, EU AI Act, ECCTA
- Per-framework control mappings + evidence-pack export

ACTIONS

Export a signed compliance pack; drill into a control's evidence.

Leads to → The compliance/ document pack

What it is. Multi-tenant administration.



KEY ELEMENTS


- Tenant list with tier, region, status, isolation settings
- Per-tenant monthly cost and alert volume

ACTIONS

Create / suspend a tenant, set data-residency region.

Leads to → Users (RBAC)

What it is. Internal readiness check — confirms the seed populates every page and acceptance criteria pass.



KEY ELEMENTS


- Green-tick matrix of pages × seed completeness
- 41-point acceptance-criteria coverage

ACTIONS

Re-run the seed check; jump to any under-seeded page.

Leads to → Any under-seeded page

What it is. Multi-channel takedown status across registrars and hosting-abuse channels.



KEY ELEMENTS


- Per-takedown channel (Cloudflare / GoDaddy / Namecheap / hosting abuse), status, reference ID
- Manual override on stalled channels

ACTIONS

Track delivery; re-route or escalate a stalled takedown.

Leads to → Audit Log

What it is. The human-in-the-loop approval surface (reviewer ≠ analyst enforced).



KEY ELEMENTS

- Pending takedowns/verdicts with the drafted packet and full verdict trace
- Second-authoriser step for egress actions (canary creds / honeypot)

ACTIONS

Approve or deny; provide a second authorisation where required.

Leads to → Takedowns (on approval) · Audit Log

Part 2 · Screen flows

A — Daily triage (the core loop)

Dashboard → Open Alerts KPI → Triage Queue (filter, pick an alert) → Alert Detail → Evidence tab (screenshots, DOM, multi-region cloaking) → Verdict tab (confirm/deny) → Takedown tab (draft) → Review Queue (HITL approve) → Takedowns (track) → Audit Log (hash-chained) → back to the queue.

B — Brand onboarding

Brand Management → Add brand / launch wizard → enter brand + regions + thresholds → wizard provisions discovery (A1-A10) → Dashboard begins to populate alerts → Brand Detail.

C — Deepfake review

Dashboard Deepfakes KPI → Deepfake Feed → open a clip (face / voiceprint / lip-sync / C2PA scores) → confirm/deny → chain-of-custody pack → Alert Detail (if tied to an alert).

D — Executive protection

Exec Protection → pick an executive → review likeness / voice findings → record lawful basis → action a redaction recommendation → escalate → Alert Detail / Audit Log.

E — Campaign investigation

Triage / Dashboard → Alert Detail → Cluster tab → Campaign Clusters graph → pivot across kit / hosting / actor nodes → open related Alert Details.

F — Agent governance & kill-switch

Admin › Agent Control → inspect agent status + budgets → (incident) press the Global Kill-Switch → all agents halt ≤ 5 s → Audit Log confirms with a hash-chained entry.

G — Cost & compliance review

Cost Attribution → review per-tenant Bright Data + LLM spend, set budgets → Admin › Compliance → check framework status → export signed evidence pack.

H — Audit investigation

Audit Log → NL search (“reviewer overrides last week”) → open a result → Alert Detail (the subject) → verify hash-chain integrity.

Part 3 · Roles, states & exports

Roles & visibility

Viewer	Read dashboards and settings.
---------------	-------------------------------

Auditor	+ Read the audit log and compliance evidence.
Reviewer	+ Approve / deny HITL actions; reviewer $\neq$ analyst is enforced.
Analyst	+ Triage, investigate, request takedowns, onboard brands.
Admin	+ Manage users, agents, cost, tenants, integrations.
Owner	+ Manage tenants and billing.

Forbidden actions are hidden; a forbidden deep link shows the **Forbidden (403)** screen.

Page states & exports. Every data screen renders one of four states — **Populated**, **Empty**, **Loading** (skeletons), **Error** (with retry). Tables export CSV / NDJSON; alerts export a signed PDF; Compliance exports a signed evidence pack; deepfake evidence exports a chain-of-custody bundle.