

SpoofVane — User Guide

Task-based: each thing you want to do, with the steps and the screen side by side.

v0.5.0 · 29 May 2026 · SOC analysts, reviewers, brand admins, auditors, administrators

About this guide

This guide is organised by **task**. For each one you get: who can do it, what you need first, numbered steps, **the screen shown beside the steps** with matching ● markers, the result, and what to do next.

Your role decides what you can do

Viewer	Read dashboards and settings.
Auditor	+ Read audit & compliance.
Reviewer	+ Approve/deny HITL actions (≠ analyst).
Analyst	+ Triage, investigate, request takedowns, onboard brands.
Admin	+ Manage users, agents, cost, tenants.
Owner	+ Manage tenants and billing.

The golden rule: every takedown and sensitive action needs a second person's approval — you propose, a reviewer approves — and everything is recorded in a tamper-evident audit log.

Contents — what do you want to do?

GETTING STARTED

T1 Sign in and find your way around

SET UP PROTECTION

T2 Onboard a brand to monitor

T3 Protect an executive from deepfakes

INVESTIGATE

T4 Triage the alert queue

T5 Review an alert and confirm the verdict

T6 Review a suspected deepfake

T7 See the campaign behind a site

ACT

T8 Approve or deny a takedown (reviewer)

T9 Track takedown delivery

OVERSEE

T10 Search the audit trail

T11 Watch Bright Data spend

ADMINISTER

T12 Manage agents and the kill-switch

Getting started

T1 Sign in and find your way around


Access SpoofVane and learn the layout. Who: All roles

BEFORE YOU START

- An account from your administrator
- SSO, or email + password
- An authenticator app if MFA is on

STEPS

- 1 Go to your SpoofVane URL — the Login screen appears.
- 2 Click SSO (OIDC) or SAML, or enter email + password.
- 3 Enter the 6-digit code if prompted.
- 4 You land on the Dashboard; the left rail is your nav.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

You're signed in on the Dashboard with today's KPIs and live pipeline.

→ NEXT

Analysts go to **Triage** (T4); to add a brand first see **Onboard a brand** (T2).

Behind the scenes: Identity (G1) handles SSO/MFA; RBAC (G2) controls what you see.

Set up protection

T2 Onboard a brand to monitor


Register a brand so discovery starts finding impersonation. Who: Analyst, Admin, Owner

BEFORE YOU START

- Signed in
- The brand's canonical login URL
- (Optional) a reference logo

STEPS

- 1 Click Brands in the left rail.
- 2 Click + New brand (wizard).
- 3 Brand name + login URL; accept the suggested keywords, TLDs and regions.
- 4 Pick the escalation score threshold (default 0.65).
- 5 You're taken to the Brand detail screen.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

The brand is live; within minutes the Triage queue fills with suspects.

→ NEXT

Watch results in **Triage** (T4).

Behind the scenes: Brand Wizard is H3; saving activates Discovery (A1–A10) over Bright Data SERP, Proxies & Datasets.

T3 Protect an executive from deepfakes


Monitor an executive's likeness and voice. Who: Analyst (with lawful basis), Admin

BEFORE YOU START

- A brand is onboarded
- A named lawful basis
- (Recommended) a reference voice clip

STEPS

- 1 Click Exec protection in the rail.
- 2 Enter the name and platforms to watch.
- 3 Required — monitoring will not start until you enter it.
- 4 Optionally upload a clean voice sample for cloned-voice matching.
- 5 Confirm to begin scanning.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

Impersonating handles, likeness/voice assets and deepfake media are reported with a risk score.

→ NEXT

Suspect media appears in **Deepfakes** (T6).

Behind the scenes: Exec surface is H5 (sensitive → governance gate E8, \$V8-5); voiceprint C9; deepfake C10.

Investigate

T4 Triage the alert queue

Work through suspects and pick what needs action. Who: Analyst

BEFORE YOU START

- A brand has produced suspects

STEPS

- 1 Each row is a suspect: verdict pill, score, source, SLA timer.
- 2 Red phish pills and short SLA timers come first.
- 3 Click a row to open Alert detail.

VERDICT	URL	SCORE	SOURCE	SLA
phish	acme-secure-login.top	0.91	SERP	2m
suspicious	acmebank-verify.xyz	0.58	OpenPhish	14m
phish	my-acme.cc	0.87	Registrar	5m
benign	acme-fanblog.net	0.12	CertStream	1h

The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

The most urgent suspect is open, ready for evidence review.

→ NEXT

Review the evidence — **Review an alert** (T5).

Behind the scenes: Ordering uses the calibrated probability C1 (not a raw sum), from Scoring C1-C8.

T5 Review an alert and confirm the verdict


Decide if a suspect is genuinely malicious using the full evidence. Who: Analyst

BEFORE YOU START

- A suspect open in Alert detail (T4)

STEPS

- 1 Verdict, calibrated score, brand, first-seen.
- 2 Open Evidence, Multi-region (cloaking?) and Kit (which kit matched).
- 3 See the Claude+GPT+Gemini ensemble; dissent discounts confidence & flags for review.
- 4 MITRE techniques, Cluster links, Deepfake signals.
- 5 Click Request takedown, or mark benign to dismiss.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

Your decision is recorded; a takedown is queued for approval, a dismissal closes the alert.

→ NEXT

A reviewer approves the takedown next — **Approve a takedown (T8)**.

Behind the scenes: Ensemble D1-D5, dissent D5, MITRE D7, clustering C8; your call feeds active-learning D8/E6.

T6 Review a suspected deepfake


Confirm or dismiss AI-generated impersonation media. Who: Analyst

BEFORE YOU START

- A deepfake suspect exists (from T3 or discovery)

STEPS

- 1 Each item shows the multimodal score.
- 2 Face distance, lip-sync, voiceprint match, C2PA provenance — valid provenance lowers suspicion.
- 3 Accept or override the multimodal verdict.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

The media is classified with per-signal reasoning kept for evidence.

→ NEXT

Request a takedown (→ T8); legal export carries a chain-of-custody pack.

Behind the scenes: Voiceprint C9, ensemble C10, fusion D6, legal custody pack G10.

T7 See the campaign behind a site


Find whether a suspect is part of a larger campaign. Who: Analyst

BEFORE YOU START

- Two or more related suspects exist

STEPS

- 1 See the threat-network graph.
- 2 Nodes connect by shared ASN, registrar, kit, favicon or TLS fingerprint.
- 3 Click any node to open that site's Alert detail.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

You see the whole campaign and can act on it as a group.

→ NEXT

Take the sites down individually (each via T5 → T8).

Behind the scenes: Clustering C8 (union-find); graph maintained by cluster agent E5.

T8 Approve or deny a takedown (reviewer)


Provide the mandatory human approval before submission. Who: Reviewer (≠ the analyst)

BEFORE YOU START

- An analyst requested a takedown (T5)
- You hold the Reviewer role

STEPS

- 1 Pending actions are listed.
- 2 Review the suspect, verdict and drafted packet.
- 3 Approve to authorise; deny with a reason.
- 4 Egress actions (canary creds / honeypot) need a second approver.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

On approval the takedown is submitted and a reference ID returns; all audited.

→ NEXT

Track delivery in **Takedowns** (T9).

Behind the scenes: HITL G4 (reviewer ≠ analyst); two-person rule via governance E8 (§V8-5).

T9 Track takedown delivery

Monitor where each approved takedown sits in the pipeline. Who: Analyst, Reviewer

BEFORE YOU START

- A takedown was approved (T8)

STEPS

- 1 See URL, channel, status and reference ID per row.
- 2 Registrar or hosting-abuse channel and its state.
- 3 Re-route or escalate manually.

URL	CHANNEL	STATUS	REF
acme-secure.top	Cloudflare	submitted	CF-6d8f
my-acme.cc	Hetzner abuse	HITL pending	—
acme-verify.xyz	Namecheap	acknowledged	NC-44a1

The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

Live visibility of every takedown and its reference.

→ NEXT

Closed alerts remain in **Audit** (T10) permanently.

Behind the scenes: Registrar adapters F1-F3, hosting-abuse F4; alerts also fan out to SIEM/SOAR via F5.

T10 Search the audit trail


Find who did what, when — for investigation or audit. Who: Auditor, Reviewer, Admin

BEFORE YOU START

- A role with audit access

STEPS

- 1 The hash-chained log opens.
- 2 Type e.g. 'analyst takedown actions for tenant acme'.
- 3 Actor, action, outcome and a hash linking each entry.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

A filtered, tamper-evident view of the relevant actions.

→ NEXT

For reporting, export an auditor pack from **Compliance** (P14).

Behind the scenes: NL search H2; hash-chained audit G3 self-verifies and detects tampering.

T11 Watch Bright Data spend


Keep data-collection cost within the envelope. Who: Admin, Owner

BEFORE YOU START

- Admin or Owner role

STEPS

- 1 See spend vs the tier cap.
- 2 Current spend, envelope, and SLM cheap-path share.
- 3 Configure a budget alert threshold.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

You can see and control cost before it exceeds the envelope.

→ NEXT

If high, review active sources in Brand detail (P04).

Behind the scenes: Cost tracking G9; every BD call records a cost row; SLM-first routing D4 keeps costs down.

T12 Manage agents and the kill-switch


Govern autonomous agents and halt them instantly if needed. Who: Admin

BEFORE YOU START

- Admin role

STEPS

- 1 See each agent, its state and blast-radius budget.
- 2 Adjust the cost/risk envelope per agent.
- 3 Global to stop everything, or per-tenant to scope it.



The numbered markers ● on the screen match the steps.

✓ RESULT — WHAT HAPPENS

Agents run within budget; firing the kill-switch halts them all (audited).

→ NEXT

Resume agents from the same screen when safe.

Behind the scenes: Kill-switch E7, governance E8, hash-chained agent audit E9.

